

# Entrepôts de Données et Big-Data

# Équipe Pédagogique

- Anne-Muriel Arigon - UM, LIRMM - Responsable UE  
([chifolleau@lirmm.fr](mailto:chifolleau@lirmm.fr))
- Federico Ulliana - INRIA, UM  
([federico.ulliana@inria.fr](mailto:federico.ulliana@inria.fr))

Special Guest :

- Christophe Menichetti - HPE (Hewlett Packard), IBM

# Objectif du Module

*vers le*

## **BIG-DATA ANALYTICS**

*Plus concrètement, présenter*

*1. les techniques de modélisation*

*2. les systèmes et technologies*

*conçues pour l'analyse des données massives*

# Programme

1. Rappels
2. Entrepôts de Données
3. Hadoop/Map-Reduce
4. **Séminaires** : Défis, Cas d'Utilisation, Solutions Big-Data Analytics & IA/ML



# Database Evolution History

1960s

First Computerized Database Models



1970s

The Dawn of the Database

- The relational model and its language SQL emerge
- The disruptive model causes the demise of other models



1980s

An Industry Develops

- SQL becomes the de-facto standard
- Commercial offerings from IBM, Oracle grow market
- Other data models enter the scene, without much traction



1990s

Technology Shifts

- Data explodes with the Internet age
- Single server SQL databases run into resource problems
- Business Intelligence and Analytics move out of transactional databases



2000s

New Players Emerge

- Data variety, velocity and volume increase
- New analytics SQL databases are introduced
- NoSQL databases fill the gap for processing unstructured data
- Hadoop gains traction for analyzing petabytes of data

Today

Databases Adapt and Evolve

- Businesses require real-time analytics on operational data
- Scale-up SQL proves too costly, but scale-out removes resource constraint
- Scale-out provides real-time analytics with high volume transactions
- Google and Cloudera are pioneers in this space

The Future

Businesses Advance with Database Innovations

- Single node SQL gets replaced by scale-out SQL
- Data warehouse type analytics will become available in real-time databases
- Businesses gain a significant edge and increased agility

Winning Database Platforms



Source: [Robin Purohit](#)

# MCC 40%TD + 60%Examen

- TD obligatoires
  - Rappels
  - Entrepôts D. 1
- TD facultatifs
  - Optimisation
  - Entrepôts D. 2
  - Map/Reduce
- Mini-projet
  - Rendu intermédiaire (document)
  - Rendu final (document, code, transparents)
- **Présence obligatoire** séminaires C. Menichetti

# Divers

- TD
  - salles 36.207 (F. Ulliana) + 5.130 (Anne-Muriel Arigon/Julie Cailler)
  - éléments de correction dispensés en cours
  - binômes ou monômes (sans exceptions)
- Consultez la page Moodle dédiée à l'UE
- Access comptes Oracle *possible à distance* (SSH)
  - instructions sur Moodle
- Signaler les dysfonctionnements des serveurs Oracle
  - ENT --> Assistance --> Centre de Services --> Déclarer un Incident
- Nous écrire pour tout type de problème ou question
  - [chifolleau@lirmm.fr](mailto:chifolleau@lirmm.fr), [federico.ulliana@inria.fr](mailto:federico.ulliana@inria.fr)

# Objectifs Pour Aujourd'hui

- Rappels
  - UML
  - Modèle Relationnel
  - SQL
- Stockage
  - “the database game” (clickbait)
- TD
  - Modélisation UML, SQL (facile)
  - Rendu jeudi 21/09 avant 23h59



# ***Rappels***

# Readings

*These slides should be considered simply as “pointers” to the references below.*

- [BD-G] Bases de données,  
Georges Gardarin, 5ème edition 2005  
[http://georges.gardarin.free.fr/Livre\\_BD\\_Contentu/XX-TotalBD.pdf](http://georges.gardarin.free.fr/Livre_BD_Contentu/XX-TotalBD.pdf)
- [ORA] Oracle® Database SQL Language  
Reference 11g Release 1 (11.1) - 2013  
[http://docs.oracle.com/cd/B28359\\_01/server.111/b28286.pdf](http://docs.oracle.com/cd/B28359_01/server.111/b28286.pdf)
- [UML] Prolegomenes\_uml.pdf
- [UML2] UML2 : de l'apprentissage à la pratique

# Relational Databases & UML

***NB : Assumed to be well known from L2/L3, we just recall basic topics.***

1. UML

2. Relational Model

3. SQL

# Levels of Modelling

- Conceptual Model (UML, EA, Merise)
  - defines what the system contains
- Logical Model (Relational Model, Object Model, Graph)
  - Defines data structures and rules of the system
- Physical Model (SQL, OQL, XML)
  - defines how the system has to be implemented



Peter Chang  
EA – 1976



Ted Codd  
RM – 1970



Don Chamberlin  
SQL – 1974



# **BASIC RELATIONAL MODEL THEORY**

# The Relational Model

[BD-G] chapter VI section 2.1

*Everything is a relation*

- **Person(Bob, 42, Paris)**  
(can model entities)
- **LiveTogether(Alice, Bob, Lyon, 2010)**  
(can model associations)

# Relational Schema

[BD-G] chapter VI section 2.2

- A set of relations built on a set of attributes, with well defined domains.

**Person**(Name, Age, City)

Name : String

Age: Integer

City : {Lyon, Paris}

# The Model VS. The Content

- The idea of representing data using relations is abstract, universal.
- But, as this data is originated from real world interactions (eg., trading, social), all forms of weak and strong correlations are found in it.



# Functional Dependency

[BD-G] chapter VI section 3

*A set of attributes **A** determining a set of attributes **B***

(name, surname) -----> birthday  
determine

city -----> (population, state)  
determine

# Key(s)

[BD-G] chapter VI section 3.1

*minimal set of attributes determining a whole tuple*

**LiveTogether**(Person1, Person2, City, Date)

key      Person1 , Person2 , Date      ----->      City  
   determines

(ex)      **LiveTogether**(Alice, Bob, Lyon, 2010)

Strongly recommended in systems (efficiency, integrity)

# Data dependencies were undesirable

- Except for **keys** and referential **integrity constraints**
  - beside these cases, they just bring redundancy
- Database **normalization** eliminated dependencies

# Normal Forms

Normal-Forms are guidelines for modeling.

Normal forms came after the relational model.

Their definition is motivated by design mistakes.

*So, let's find the right place for the attributes !*

# Normal Forms : 2NF

[BD-G] chapter VI section 6.2

- **2NF** : non-key attributes fully-dependent from the key

**FournisseurPiece**(Name, Article, Address, Price) **NO**

(Article --/--> Address)

↓ (decomposition)

**FournisseurPiece**(Name, Article, Price)

**Fournisseur**(Name, Address)

# Normal Forms : 3NF

[BD-G] chapter VI section 6.3

- 3NF : no dependencies between non-key attributes

**Person**( ID , Name, City, *CityPopulation* )

**NO**

ID -----> City -----> CityPopulation



(decomposition)

**Person**( ID , Name, City)      **Place**(City, *CityPopulation*)

# Normal Forms

[BD-G] chapter VI section 6.3

- 3NF respected by **most** “*transactional-database*” you will find in **any** real world company
  - it allows to fix common data redundancy problems
  - also, every schema can be normalized in 3NF
- but sometimes only 2NF ...
  - not (always) by mistake of database administrator
  - conscious choice : pay the price of redundancy (and update) to gain in performances (more later)
- 3NF = 2NF + no indirect dependencies (simplification)
- 2NF = 1NF + no partially dependent attributes (simplification)
- 1NF = ?
- Are there transactional databases which do not respect the 1NF ? (Non-first-normal-form N1NF)

# Normal Forms

[BD-G] chapter VI section 6.3

- 3NF respected by **most** “*transactional-database*” you will find in **any** real world company
  - it allows to fix common data redundancy problems
  - also, every schema can be normalized in 3NF
- but sometimes only 2NF ...
  - not (always) by mistake of database administrator
  - conscious choice : pay the price of redundancy (and update) to gain in performances (more later)
- 3NF = 2NF + no indirect dependencies (simplification)
- 2NF = 1NF + no partially dependent attributes (simplification)
- 1NF = only atomic values in a tuple
- Are there transactional databases which do not respect the 1NF ? (Non-first-normal-form N1NF)
  - Yes ! Object-oriented, XML, many in the NoSQL family (using JSON as data model)



# Normal Forms : Remarks

- 2NF & 3NF respected in practically any information system using a relational database
- Stronger normal forms (BCNF, 4NF, 5NF) are less employed (avoid rarer mistakes; not always achievable)
- Exceptions to normalizations are tolerated to save joins (at the price of redundancy)
  - **we will see this for datawarehouses**



# SQL : SURVIVAL KIT

# SQL

- Structured Query Language
- Declarative (logical) Language : tell what you want from relations, not what to do with them.
  - This is the main difference with C and Java, *not only* the fact that we deal with data.
- In SQL terminology, a relation is called “table”.

# SQL

- DDL (Data Definition Language)
  - CREATE/ALTER structures (table, view, index)
- DML (Data Manipulation Language)
  - UPDATE/INSERT/DELETE content
- DQL (Data Query Language)
  - SELECT data

# SQL

- The purpose of the following slides is just to recall you the basic syntax.
- Check the documentation for more details or advanced features !

# Create Table

[BD-G] chapter VII section 2.1 and [ORA] section 16-6

```
CREATE TABLE Employee (  
  
    id    NUMBER,  
    name  VARCHAR2(50),  
    birthday DATE  
  
)
```

# Datatypes

[ORA] section 2-6

Why do we need datatypes ?

- To associate fixed set of properties to attributes.
- This improves the database:
  - coherence : type-checking operations
    - cannot sum two strings
  - efficiency : a datatype has its own best storage
    - BLOB vs integers

# ORACLE Built-in Datatypes

- Character
- Numeric
- Date/Time
- Large Object

Complete list : see [ORA] table 2-1



# Value Constraints

[ORA] section 8-4 and [BD-G] chapter VII section 2.1

- Why do we need constraints ?
- To restrict the values in a database and ensure the data integrity
  - ex : No employee without an ID

# Not Null

[ORA] section 8-8

- Prohibits a database value from being null

```
CREATE TABLE Employee (  
  
  id          NUMBER NOT NULL,  
  name        VARCHAR2(50),  
  birthday    DATE  
  
)
```

# Unique

[ORA] section 8-9

- Prohibits multiple rows from having the same value (but allows them to be null)

```
CREATE TABLE Employee (  
  
    id            NUMBER UNIQUE,  
    name          VARCHAR2(50),  
    birthday      DATE  
  
)
```

# Primary Key

[ORA] section 8-9

- Combines a NOT NULL constraint and a UNIQUE constraint in a single declaration

```
CREATE TABLE Employee (  
  
    id            NUMBER PRIMARY KEY,  
    name          VARCHAR2(50),  
    birthday      DATE  
  
)
```

# Primary Key : Multiple Attributes

[ORA] section 8-20

- Combines a NOT NULL constraint and a unique constraint in a single declaration

```
CREATE TABLE Employee (  
  id          NUMBER,  
  name        VARCHAR2(50),  
  birthday    DATE,  
  
  PRIMARY KEY (name,birthday)  
)
```

# Foreign key

*one (or more) **attributes** which correspond to the **key** of another relation*

**Employee( ID, Name, **Department\_id** )**

**Departement( Dept\_ID, Name )**



# Foreign Key

[ORA] section 8-10 and 8-21

Requires values in one table to match values in another table.

```
CREATE TABLE Employee (  
    id          NUMBER,  
    department  NUMBER  
  
    FOREIGN KEY department  
                REFERENCES Dept(dept_id)  
)
```

# Check

[ORA] section 8-10 and 8-22

Requires a value to satisfy with a specified condition

```
CREATE TABLE Employee (  
    id                NUMBER,  
    department        NUMBER,  
    office            VARCHAR2(10)  
  
    CHECK (  
        office IN  
            ('DALLAS', 'BOSTON', 'PARIS', 'TOKYO')  
    )  
)
```

Oracle does not verify mutually exclusive conditions (eg. AGE>1 AND AGE<0)



# ALTER TABLE

[BD-G] chapter VII section 6.2.4 and [ORA] section 12-2

- Add a new column
  - `ALTER TABLE Employee ADD (office VARCHAR2(20));`
- Modify an existing column
  - `ALTER TABLE Employee MODIFY (office NUMBER);`
- Define a default value for the new column
  - `ALTER TABLE Employee MODIFY office DEFAULT 'Corridor';`
- Drop a column
  - `ALTER TABLE Employee DROP (office);`

# DELETE

# TRUNCATE

# DROP

[ORA] section 17-25 and 19-62 and 18-5

removes

rows

table

rollback

DELETE



TRUNCATE



DROP



- TRUNCATE = DROP + CREATE TABLE

# INSERT

[BD-G] chapter VII section 4.1 [ORA] section 18-66 and 18-54

INSERT

INTO Employee

VALUES

( ' Bob ' ,

***TO\_DATE***(

' 03-OCT-1972 ' , ' DD-MON-YYYY ' )

)

***TO\_DATE*** converts a character/numeric to a date

[ORA] 2-49/50

# SELECT FROM

[BD-G] chapter VII section 3.1 [ORA] section 18-66 and 18-54 and 2-49

SELECT

*TO\_CHAR*(birthday, 'MM-DD-YYYY' )

FROM

Employee

*TO\_CHAR* converts a numeric to a character

[ORA] 2-287/288 and 5-292

# SELECT FROM WHERE

[BD-G] chapter VII section 3.2 [ORA] section 2-50

SELECT

name

FROM

Employee

WHERE

birthday >

*TO\_DATE*('01-10-1970', 'DD-MM-YYYY')

# JOINS

[BD-G] chapter VII section 3.3

SELECT

Employee.name, Department.name

FROM

Employee, Department

WHERE

Employee.dept = Department.id

# JOINS

*Employee*

<b>name</b>	<b>dept</b>
Alice	dep1
Bob	dep2
Eddy	dep1

*Department*

<b>id</b>	<b>name</b>
dep1	Sales
dep2	Production

*Emp\_Join\_Dep*

<b>Employee.name</b>	<b>Department.name</b>
Alice	Sales
Bob	Production
Eddy	Sales

# Group By

[BD-G] chapter VII section 3.7

```
SELECT
    dept, count(*) as N
FROM
    Employee
GROUP By
    dept
```

*Employee*

name	dept
Alice	dep1
Bob	dep2
Eddy	dep1

*Agg\_Emp*

dept	N
dep1	2
dep2	1

*Group by queries are at the essence of ANALYTICS.*



# Summing Up

- Relations
  - Functional dependencies, Normal-forms
- SQL
  - CREATE, INSERT, DELETE, SELECT, GROUP-BY



# **CONCEPTUAL MODELING WITH ENTITY-RELATIONSHIP DIAGRAMS**



*To minimize the number of graphical notations introduced across the various programs of this training, we take an alternative model.*

# **CONCEPTUAL MODELING WITH ~~ENTITY-RELATIONSHIP DIAGRAMS~~**

# **CONCEPTUAL MODELING WITH UML**

# UML (Unified Modeling Language)

- Universal **graphical** modeling language designed to model objects, associations, time events, system states
  - Main goal is to ease prototyping
- Good news : UML is a rich model and we can also use it to model data !

# UML (Unified Modeling Language)

## **graphical**

means :

- flexible
- error-prone
  - improper use of notation, wrong modelling
- no debugger !


# UML : Plan

- UML Class Diagram
  - Basic constructs that can be used to model Relational Databases in UML
  - Real Object-oriented features

# Class

[UML] section 3 [BD-G] chapter XVII section 2

- Set of elements sharing common properties
  - ex. peoples, animals, cars
- UML : draw a labelled box






# Class : Attributes

[UML] section 3 [BD-G] chapter XVII section 2

- Attributes denote the properties of class objects
  - They are usually typed
- Write the attributes below the class name




# Operations can specify

- Visibility (+ public) (- private) (# protected)
- Return-type (*optional*; can be undefined)
- Multiplicity of parameter/return-type (*optional*)

Employee
name : string dept : int
<b>+setDepartment( int dept [1]) : bool [1]</b>

# Instances

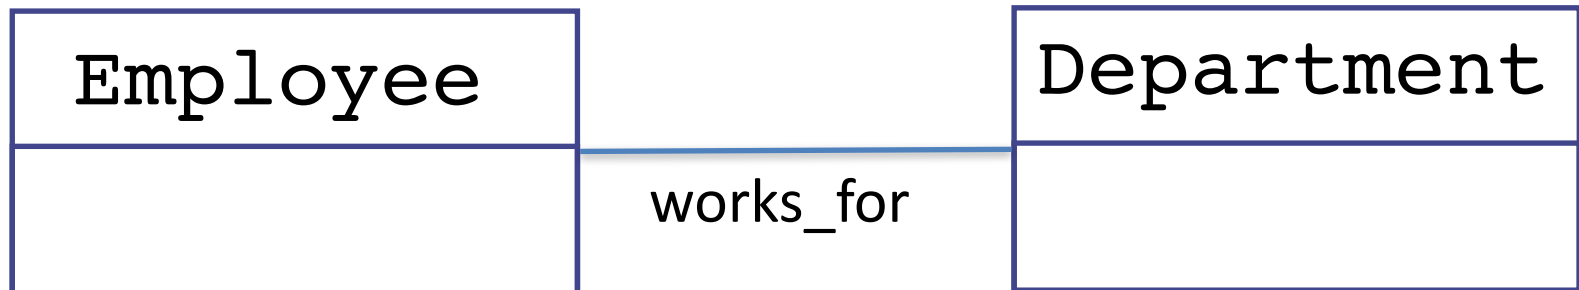
- The elements of a class
  - ex. the employee Bob
- Related to their class by a directed edge



# Binary Associations

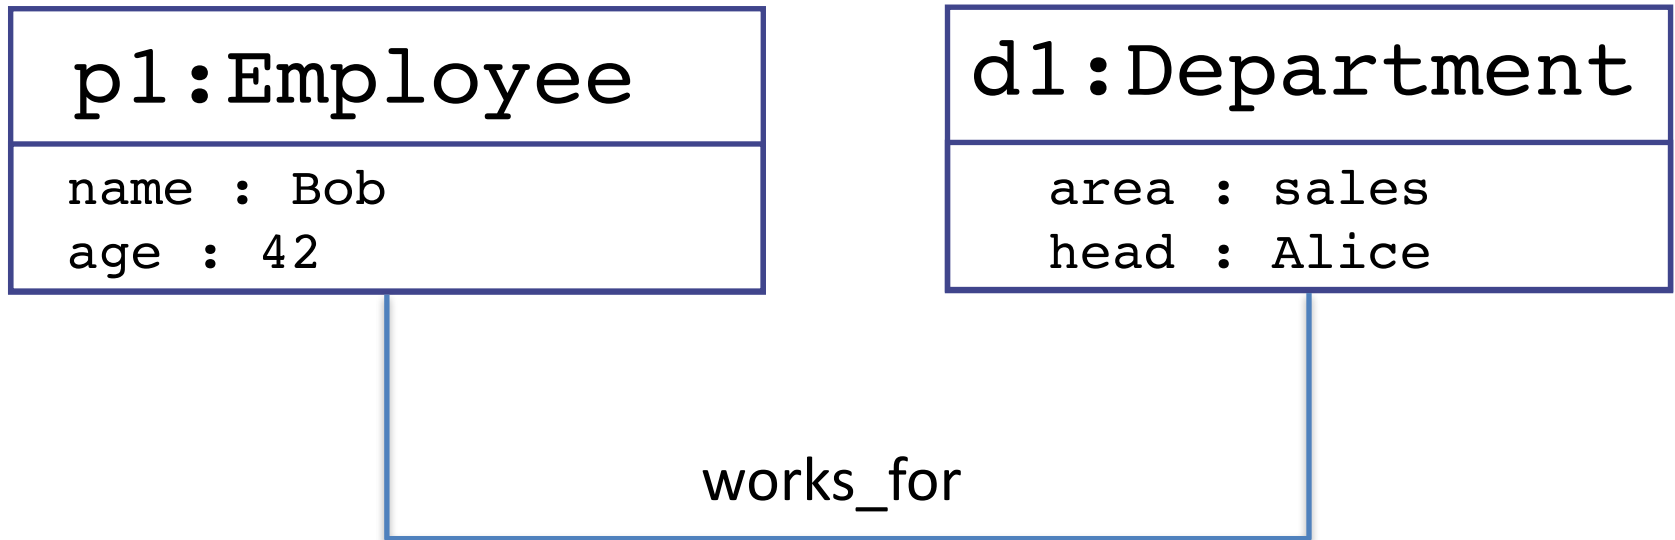
[UML] section 4 [BD-G] chapter XVII section 2

- General relationships between elements of two classes
  - ex. an employee works for a department
  - a concrete instance of association is called link
- Binary association : *undirected* edge between classes



# Associations : Links


- Relationships between instances of classes



# Associations : a tricky notation


- Any association is specified by three things
  1. Its name
  1. The cardinality constraints of the class elements
  2. The role of the classes in the association (optional)

# Associations : Name (1)



- The name is a label placed in the middle of the edge


# Associations : Cardinality (2)



- This means that one instance of A participates in the association with **N** elements of the other class




# Associations : Cardinality (2)




This means that an employee works for exactly **1** department

# Associations : Cardinality (2)



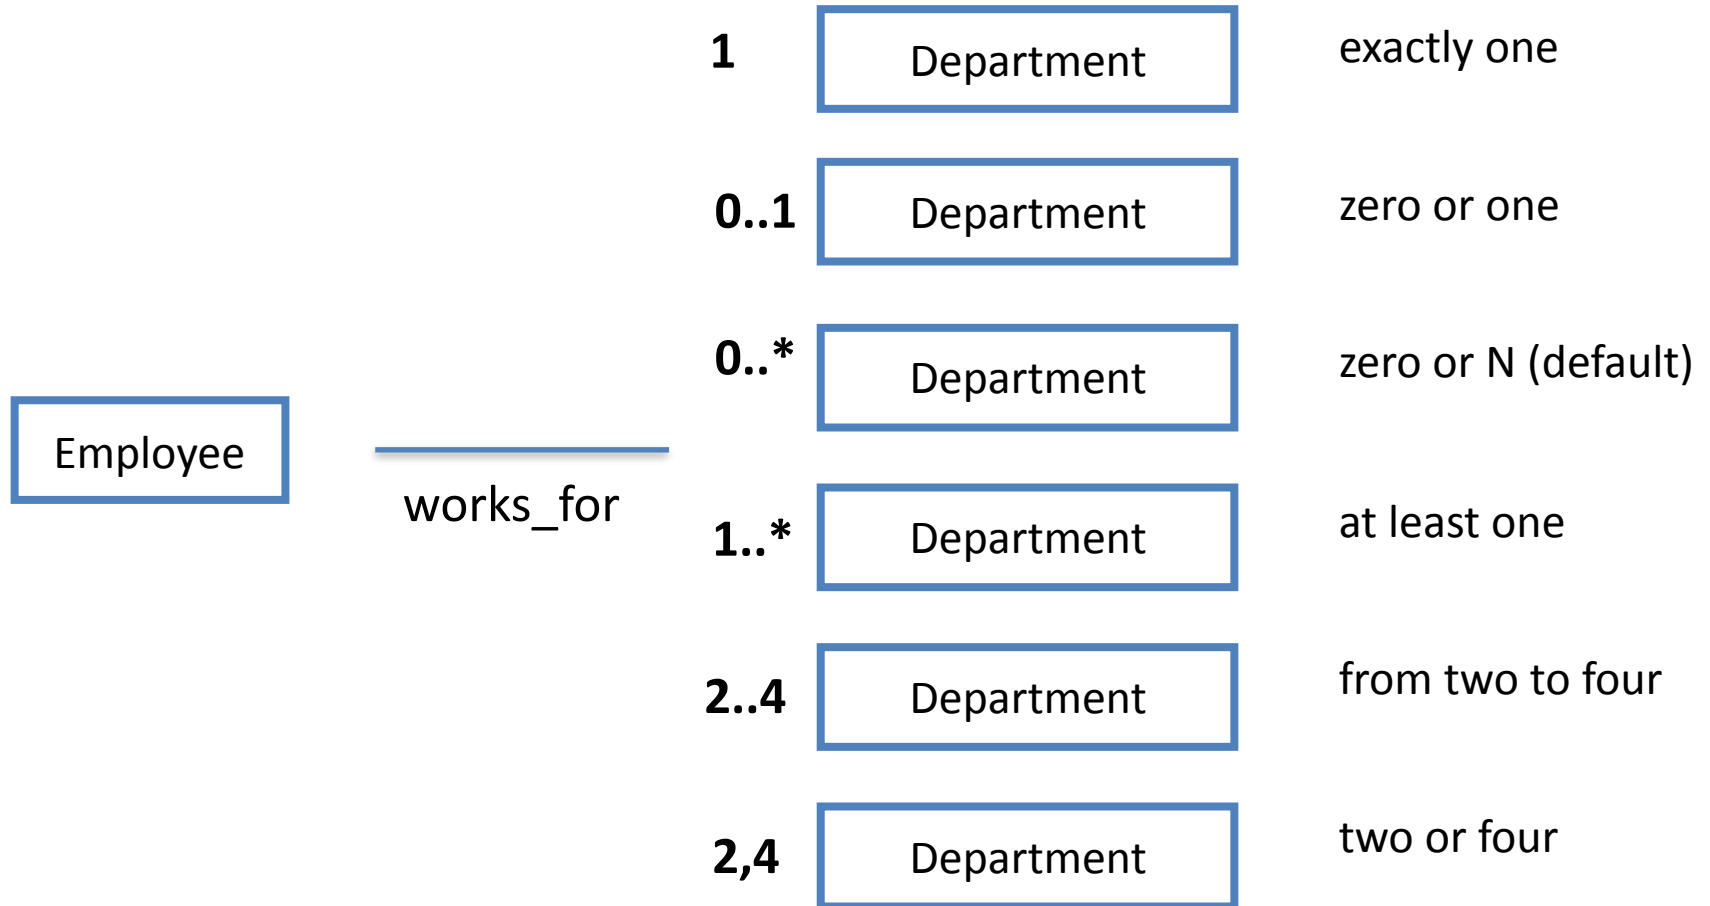
This means that an employee can work for more than **1** department (but at least one)

# Associations : Cardinality (2)




This means that a department has at least one employee, with no upper-limit.

# Cardinality Specification

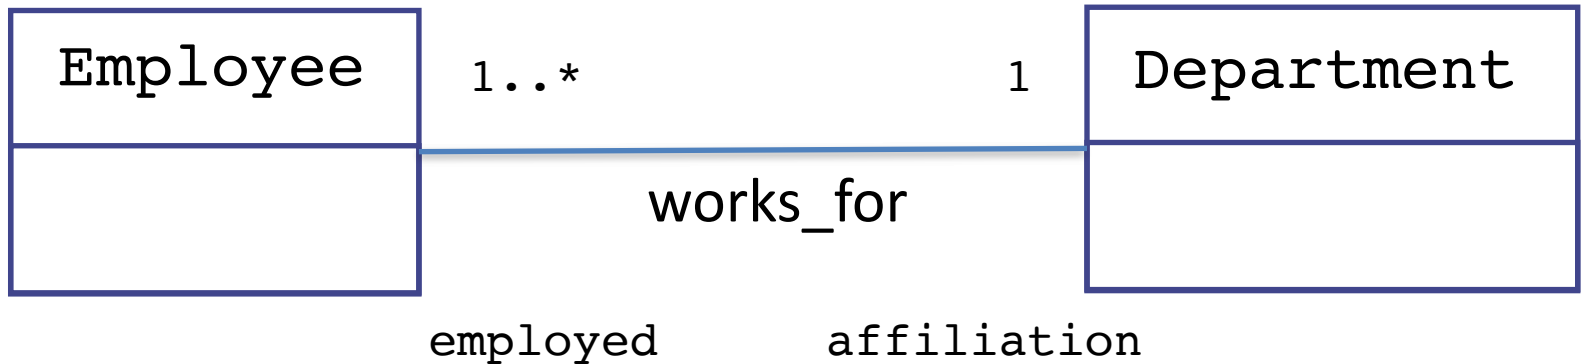


## Associations : Roles (2)

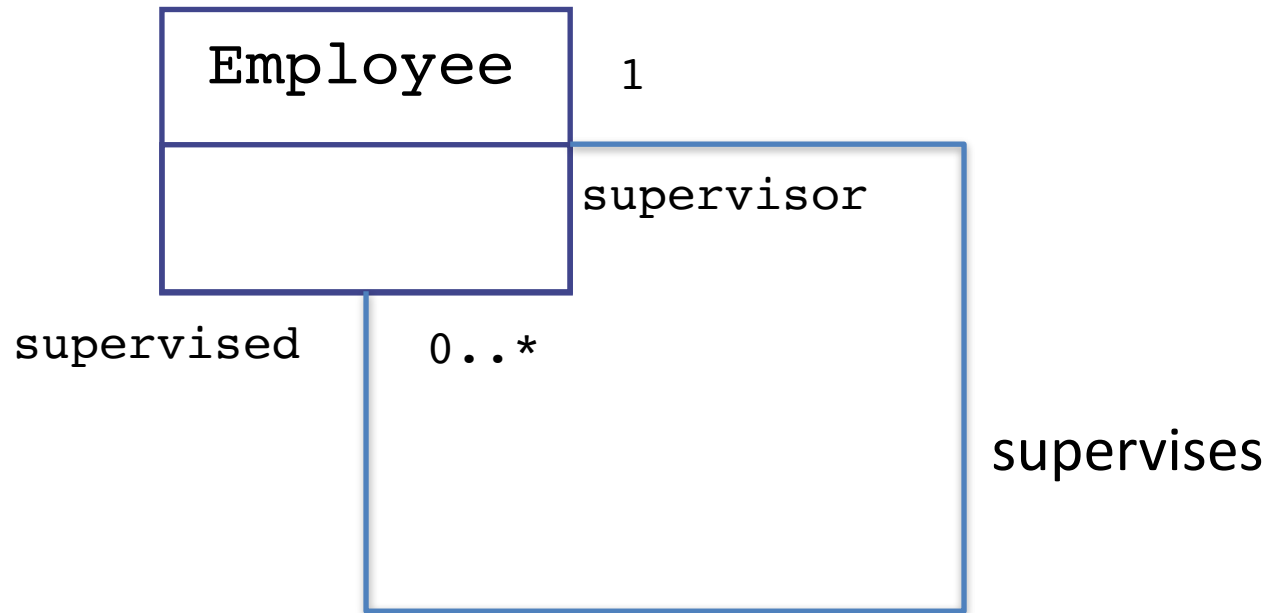


- This means that in the association an instance of A plays **role<sub>A</sub>**

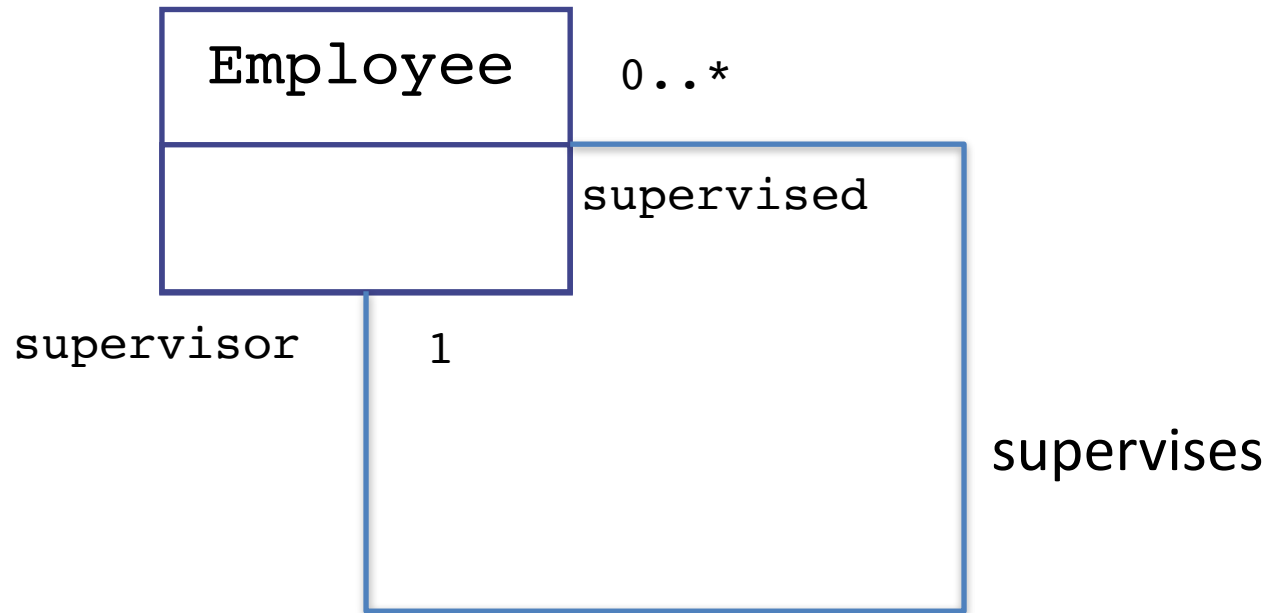
# Putting everything together



# Reflexive Associations



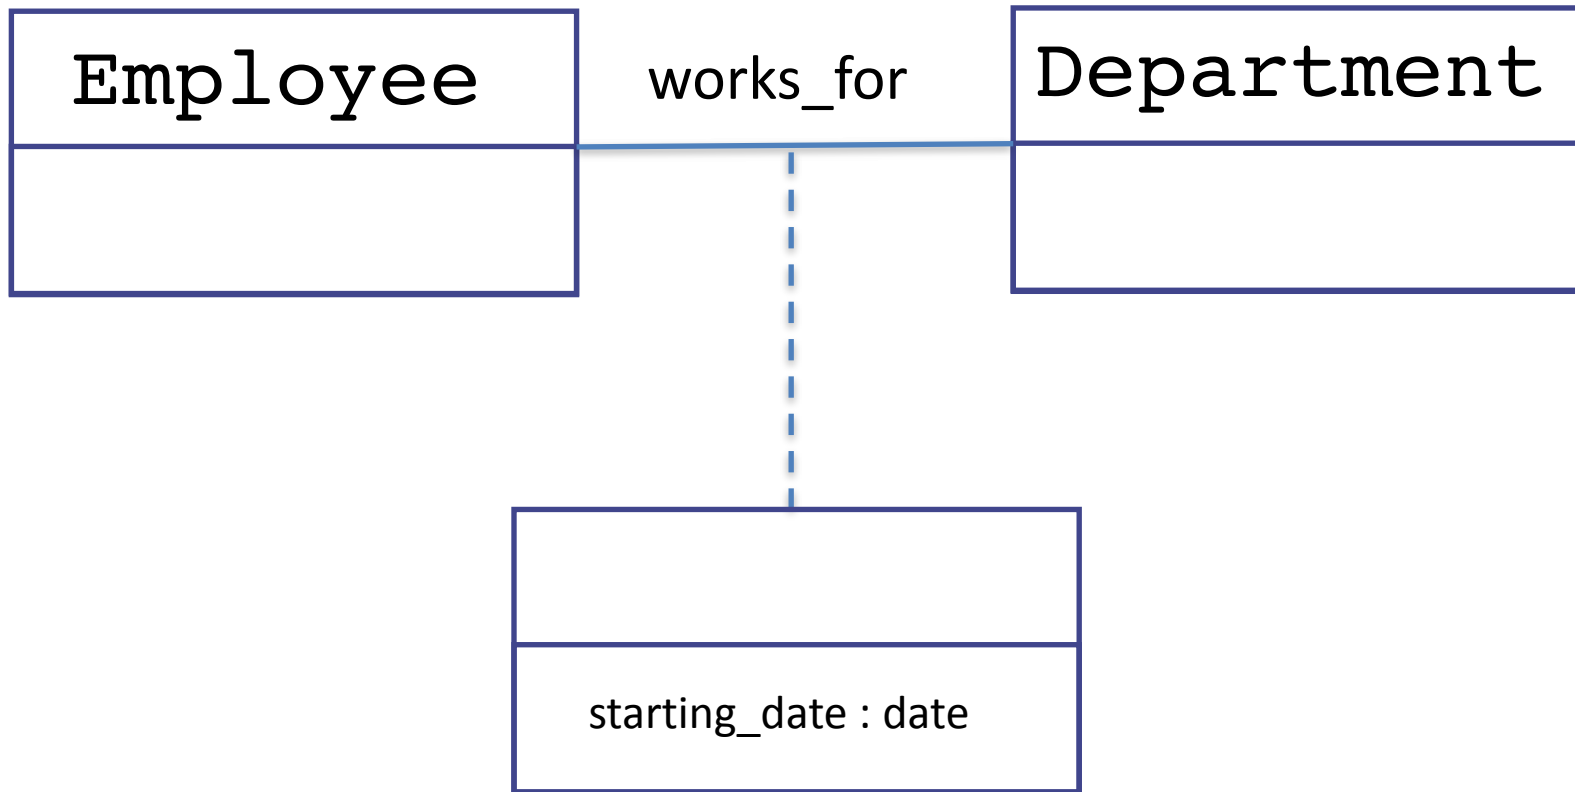
# Equivalent Formulation





# Notation for Attributes in Associations

[UML] section 4 [BD-G] chapter XVII section 2



# Associations, Associations, Associations...

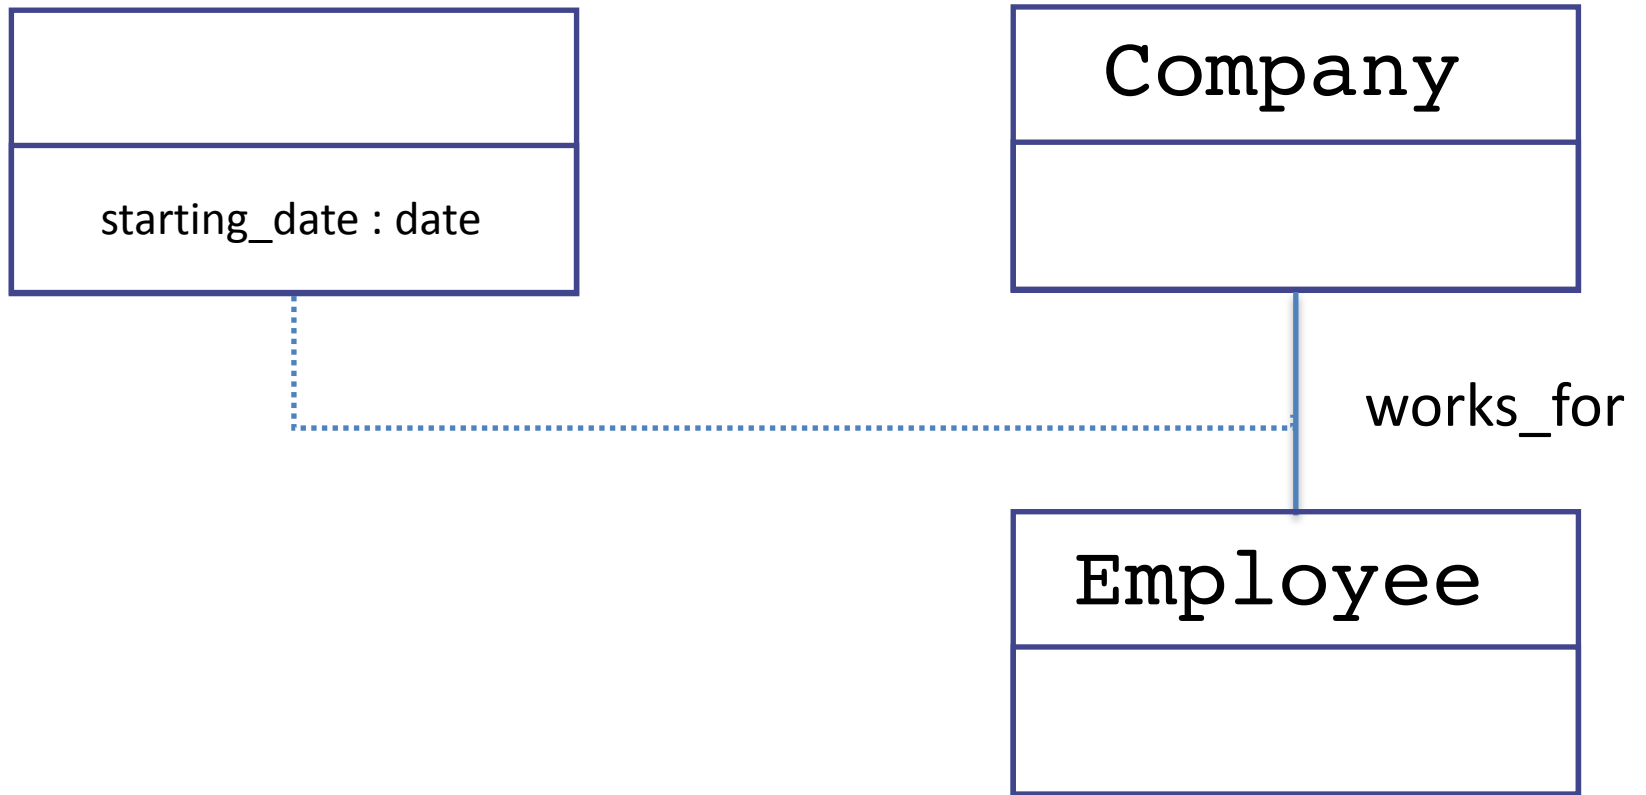
Often, a relationship between two entities combines:

- *association* features
  - the fact that two or more things are linked
- *representation* features
  - the details about this association
- Ex: a working contract can be seen
  - as a relationship  
(an association between an employee and a company)
  - as an entity (representation of a legal concept)

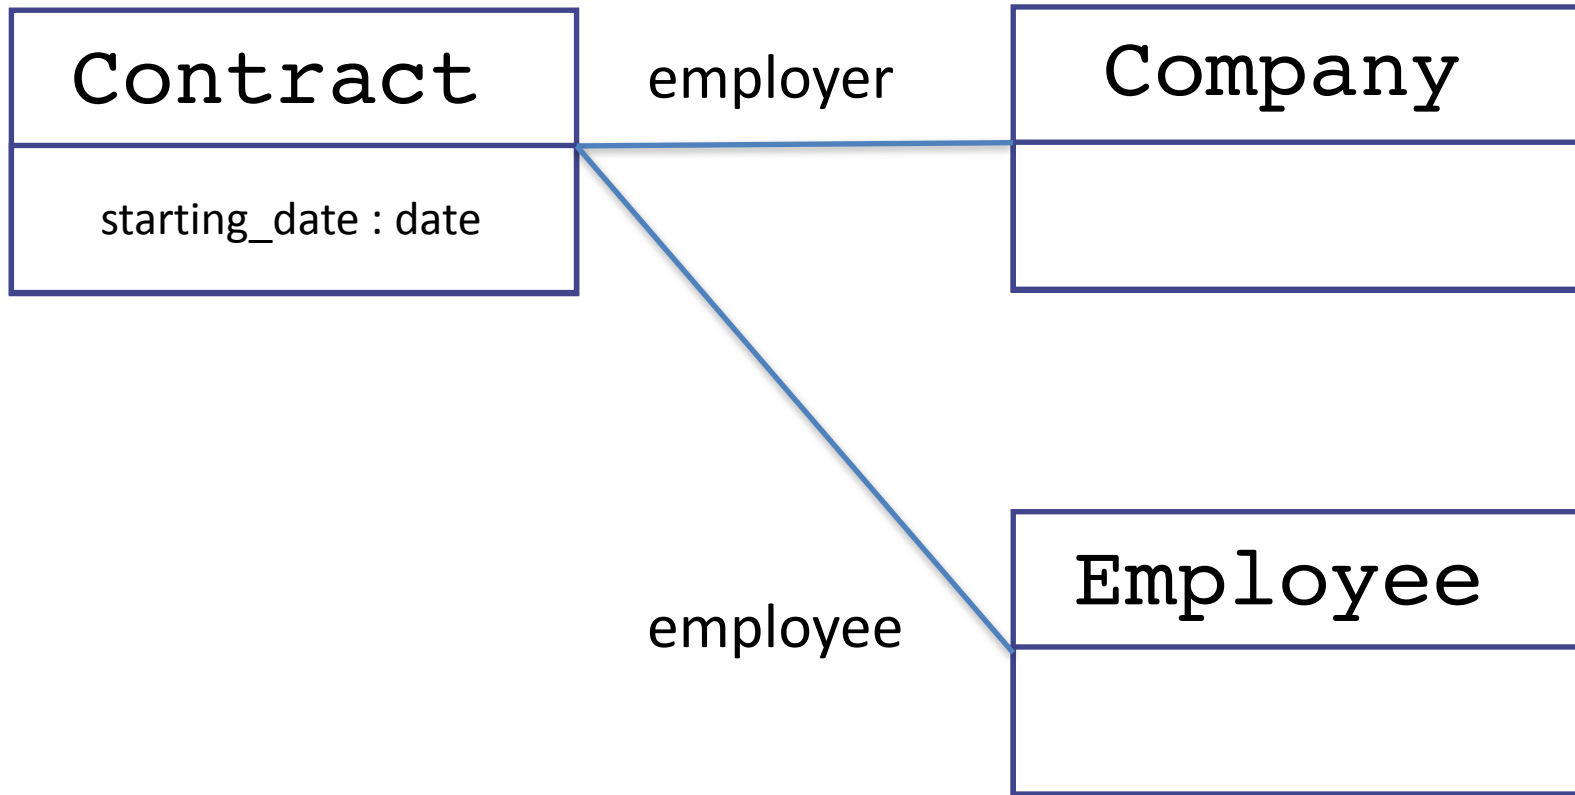
This duality is the source of most design problems !

- Recognize your own modelling choices !

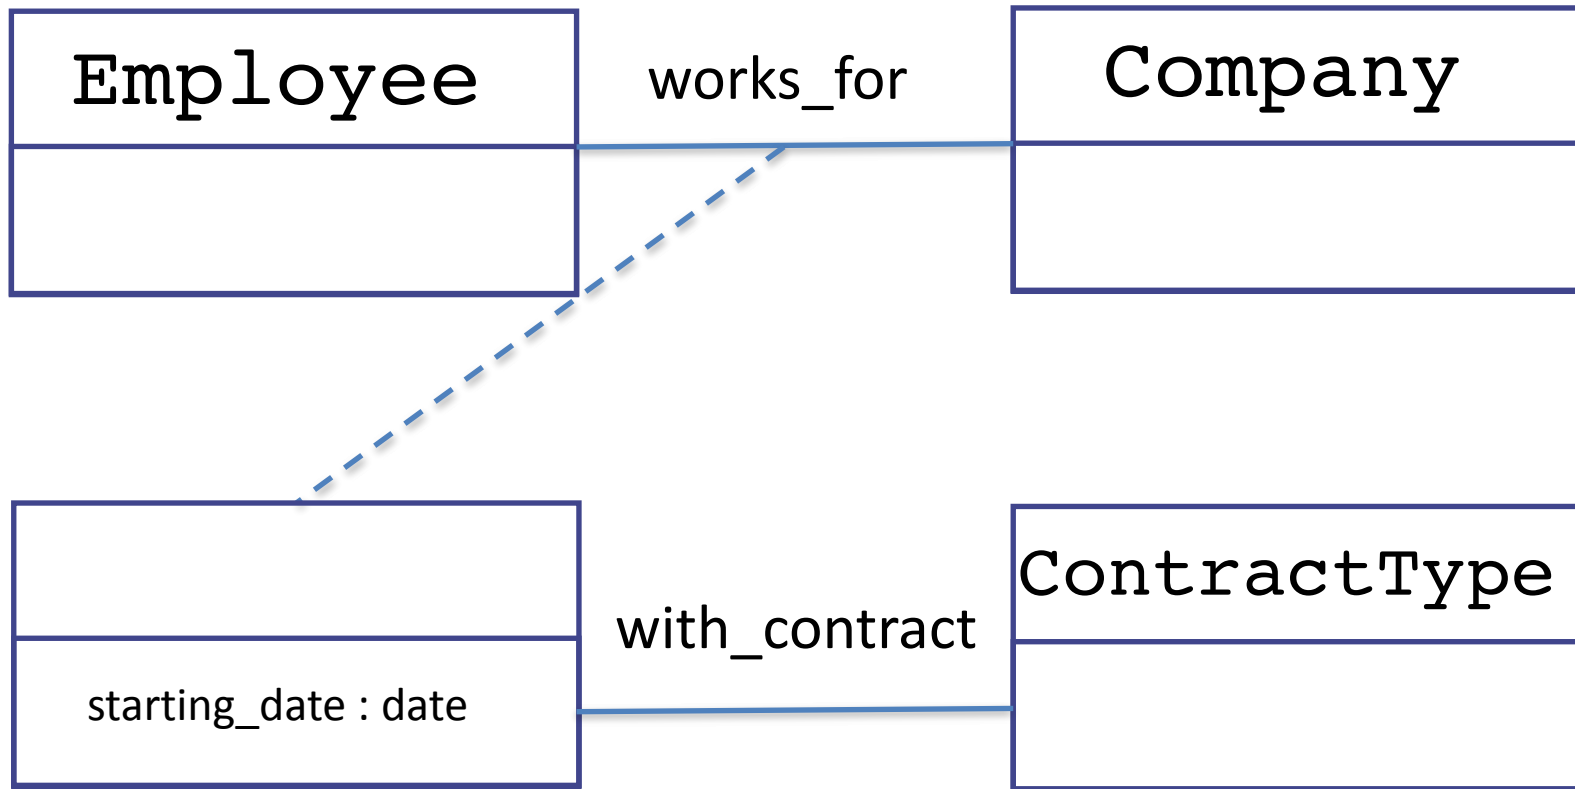
# Choice 1



# Choice 2




# Associations Participating in Other Associations



# Notation for 3-ary Associations

[UML] section 4.1 [UML2] section [3.3.4]



# Summing Up

UML for Relational Databases


[UML] section 4 [BD-G] chapter XVII section 2

- Basic UML
  - Classes, binary associations, n-ary associations
  - Can model relational databases
- The construct of the language we have seen so far are enough to model relational databases

# SubClass

[UML] section 5 [BD-G] chapter XVII section 2.2


- A subset of the instances of a Class
  - ex. every employee is a person
- A subclass inherits all superclass properties or redefines them.





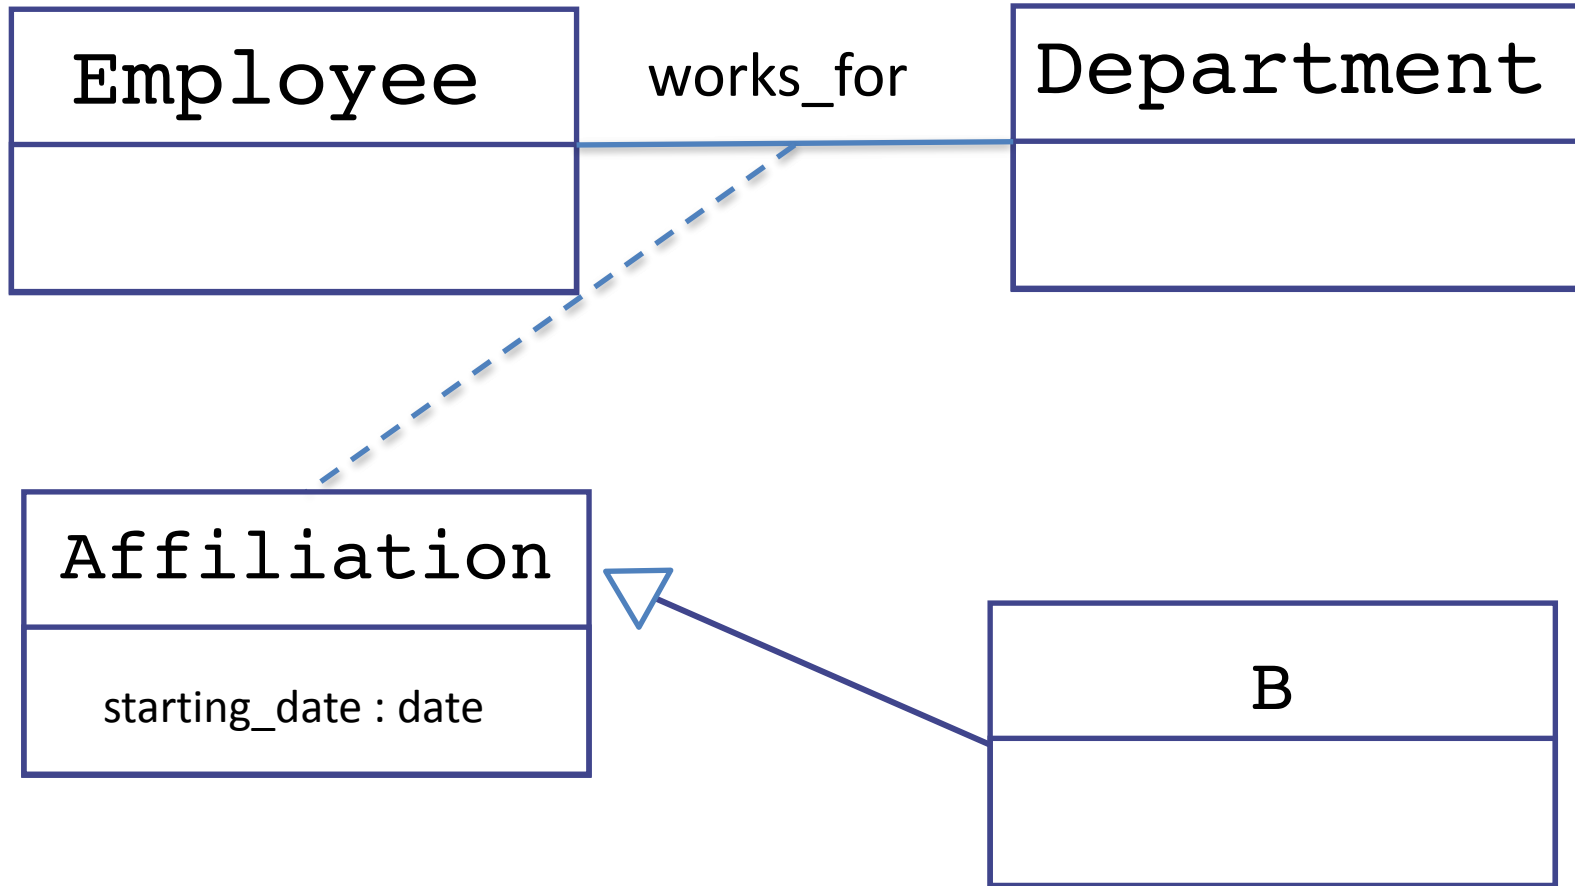
# SubClass = Generalization/Specialization

[UML] section 5

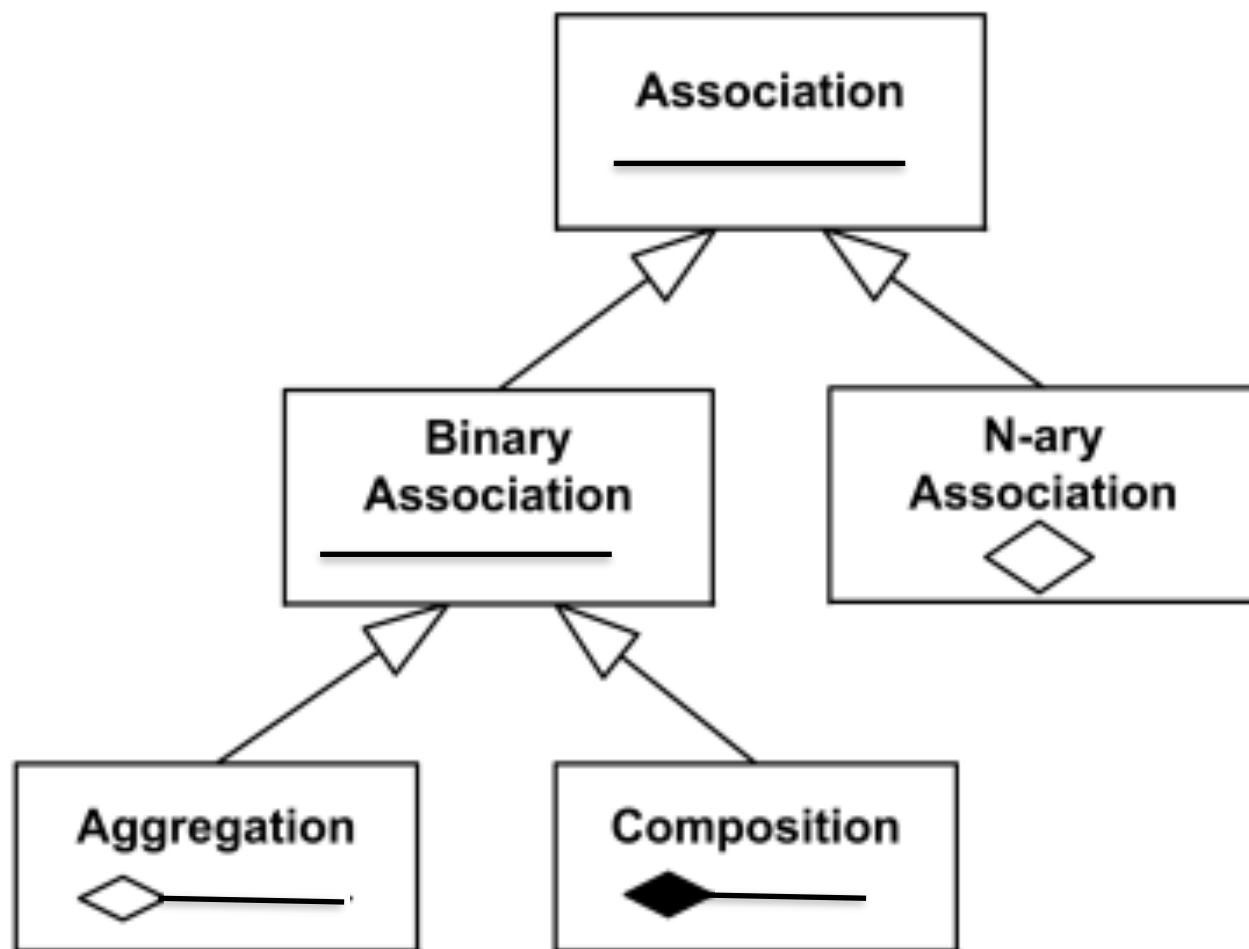


# Association modeled by a Class

[UML] section 4.3



- Consequence : such class can have proper attributes, operations and also other associations




# Part-Of (binary) Association

- “A forest is made of trees”

# Aggregation

[BD-G] chapter XVII section 2.2

- Consequence 1 : a tree can exist even without a forest




# Part-Of (binary) Association

- “A folder is made of files”

# Composition

[BD-G] chapter XVII section 2.2


- Consequence 1 : a file cannot exists without a folder



# Aggregation


[BD-G] chapter XVII section 2.2

- Consequence 1 : a tree can exist even without a forest



# Composition


- Consequence 1 : a file cannot exists without a folder






# Aggregation

- Consequence 2 : a tree can be part of both a forest and a park

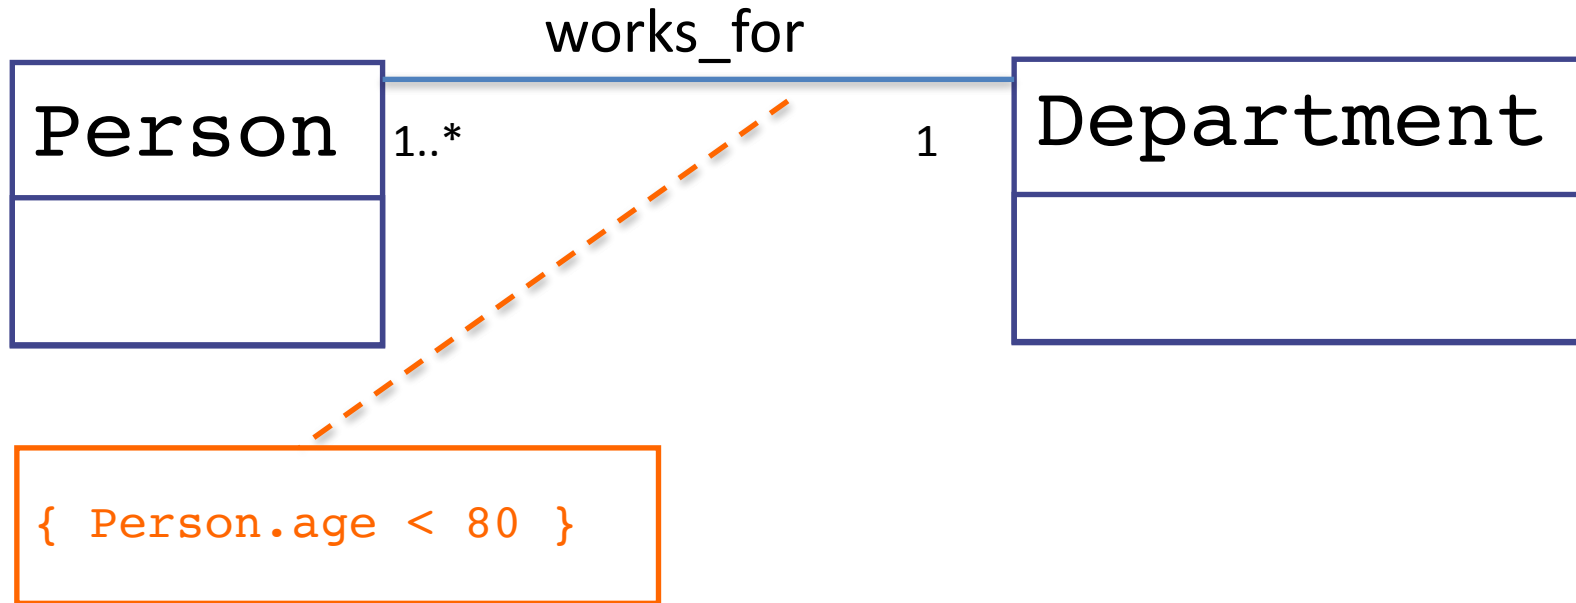


# Composition

- Consequence 1 : a file cannot exists without a folder



# Constraints on Associations



- Consequence : only people whose age is less than 80 can participate in the association

# Summing Up

- Instances/Links/Operations
- Subclasses
- Aggregation and Composition